


EECS 836: Machine Learning

Zijun Yao

Assistant Professor, EECS Department
The University of Kansas

Language Model Capability Comparison



How to determine whether the output is correct based on the standard answer

How about multiple-choice questions

- Massive Multitask Language Understanding (**MMLU**)

<https://arxiv.org/abs/2009.03300>

Find all c in $\mathbb{Z}_3$ such that $\mathbb{Z}_3[x]/(x^2 + c)$ is a field.

(A) 0 (B) **1** (C) 2 (D) 3

A point pole has a strength of $4\pi \times 10^{-4}$ weber. The force in newtons on a point pole of $4\pi \times 1.5 \times 10^{-4}$ weber placed at a distance of 10 cm from it will be

(A) **15 N.** (B) 20 N. (C) 7.5 N. (D) 3.75 N.

From the solubility rules, which of the following is true?

(A) All chlorides, bromides, and iodides are soluble

(B) All sulfates are soluble

(C) All hydroxides are soluble

(D) **All ammonium-containing compounds are soluble**

Different implementations of the same MMLU evaluation	MMLU (HELM)	MMLU (Harness)	MMLU (Original)
llama-65b	0.637	0.488	0.636
tiiuae/falcon-40b	0.571	0.527	0.558
llama-30b	0.583	0.457	0.584
EleutherAI/gpt-neox-20b	0.256	0.333	0.262
llama-13b	0.471	0.377	0.47
llama-7b	0.339	0.342	0.351
tiiuae/falcon-7b	0.278	0.35	0.254
togethercomputer/RedPajama-INCITE-7B-Base	0.275	0.34	0.269


<https://huggingface.co/blog/open-llm-leaderboard-mmlu>

Issues exist even for multiple-choice questions

Find all c in $\mathbb{Z}_3$ such that $\mathbb{Z}_3[x]/(x^2 + c)$ is a field.
(A) 0 (B) 1 (C) 2 (D) 3

(The correct answer is B)

Only output the choice, not any other contents



Answer is b

Don't forget AI can generate any thing


Are we testing the real capacity of math, or
just the understanding of instructions?

Through the calculation, I
would choose 1

Make machine to choose

Find all c in $\mathbb{Z}_3$ such that $\mathbb{Z}_3[x]/(x^2 + c)$ is a field.
(A) 0 (B) 1 (C) 2 (D) 3

(B is correct answer)



- ☒ Correct : In A, B, C, D answers, B has the highest score
- ☒ Wrong : Model obviously answers "3"

Make machine to choose

Move all the
correct answer to A




<https://arxiv.org/abs/2309.03882>

I like guessing A


			Move Golden to	Orig	A	B	C	D
			llama-30B	53.1	68.2 (+15.2)	54.1 (+1.1)	50.1 (-2.9)	41.2 (-11.9)
			vicuna-v1.3-33B	57.0	59.5 (+2.5)	58.6 (+1.5)	65.8 (+8.8)	44.8 (-12.3)
			falcon-40B	51.8	46.3 (-5.5)	45.2 (-6.7)	64.8 (+13.0)	47.9 (-3.9)
			falcon-inst-40B	51.5	38.3 (-13.3)	38.9 (-12.7)	55.7 (+4.1)	69.1 (+17.6)
			llama-2-70B	64.0	61.5 (-2.6)	68.6 (+4.5)	64.1 (+0.1)	62.0 (-2.1)
			gpt-3.5-turbo	67.2	65.3 (-1.9)	68.5 (+1.3)	74.2 (+6.9)	60.9 (-6.3)
Methods	MMLU							
	RStd	Acc						
Default	5.5	67.2						
a/b/c/d	6.8	67.0						
1/2/3/4	3.8	65.8						
(A)/(B)/(C)/(D)	8.1	66.5						

Types of Questions Without a Single Standard Answer

Translation



Summarization



Being different from the standard answer
doesn't necessarily mean it's wrong.

Maybe human evaluation is more accurate?

<https://lmarena.ai/>

Chatbot Arena

Model A

Please explain briefly how to increase my work efficiency.

Okay, here's a breakdown of how to boost your work efficiency, broken down into easily digestible points:
1. Prioritize & Plan:

- To-Do List:** Write everything down.
- Prioritize (Eisenhower Matrix):** Categorize tasks as:
 - Urgent & Important:** Do these *now*.
 - Important, Not Urgent:** Schedule these.
 - Urgent, Not Important:** Delegate if possible.
 - Neither:** Eliminate.
- Time Blocking:** Allocate specific time slots for specific tasks.

2. Minimize Distractions:

- Turn off Notifications:** Email, social media, phone.
- Dedicated Workspace:** If possible, have a designated work area.
- Communicate Boundaries:** Let colleagues/family know when you need focused time.

3. Work Smarter, Not Harder:

Model B

Please explain briefly how to increase my work efficiency.

Increasing work efficiency involves optimizing your time, energy, and resources to accomplish more with less effort. Here are some concise strategies to boost your efficiency:
1. Set Clear Priorities: Use techniques like the **Eisenhower Matrix** to distinguish between urgent and important tasks. Focus on **high-impact** activities that align with your goals.
2. Plan Ahead: Start your day by creating a **daily to-do list** or using a **time-blocking method** to allocate specific time slots for tasks. This helps you stay organized and focused.
3. Limit Distractions: Identify common distractions (e.g., social media, notifications) and eliminate or minimize them. Consider using apps like **Freedom** or **Focus@Will** to block distractions.
4. Use the Pomodoro Technique: Work in **25-minute focused bursts** (Pomodoros) followed by a 5-minute break. After four Pomodoros, take a longer break (15–30 minutes). This maintains energy and focus.
5. Batch Similar Tasks: Group similar tasks together (e.g., responding to emails, making phone calls) to reduce mental switching and increase flow. This is known as **task batching**.
6. Delegate and Automate: Identify tasks that can be **delegated** to others or **automated** using tools (e.g., email filters, scheduling software). This frees up time for higher-priority work.
7. Optimize Your Workspace: Keep your workspace **tidy** and **ergonomic**. A clutter-free environment reduces stress and improves focus.

A is better

B is better

Tie

Both are bad

Chat

Search RepoChat Text2Image

Enter your prompt and press ENTER

Send

Maybe human evaluation is more accurate?

Leaderboard (by April 2025)

Language Overview Price Analysis WebDev Arena Vision Text-to-Image Copilot Arena Search Arena-Hard-Auto

Total #models: 225. Total #votes: 2,871,517. Last updated: 2025-04-16.

Code to recreate leaderboard tables and plots in this [notebook](#). You can contribute your vote at [lmarena.ai](#)!

Category

Overall

Apply filter

☐ Style Control

☐ Show Deprecated

Overall Questions

#models: 225 (100%)


#votes: 2,871,517 (100%)

Rank★ (UB)	Rank (StyleCtrl)	Model	Arena Score	95% CI	Votes	Organization	License
1	1	Gemini-2.5-Pro-Exp-03-25	1439	+7/-5	9013	Google	Proprietary
2	2	ChatGPT-4o-latest (2025-03-26)	1407	+6/-6	8261	OpenAI	Proprietary
2	4	Grok-3-Preview-02-24	1402	+5/-3	14849	xAI	Proprietary
2	2	GPT-4.5-Preview	1398	+5/-6	14520	OpenAI	Proprietary
2	4	Gemini-2.5-Flash-Preview-04-17	1392	+10/-13	3325	Google	Proprietary
5	9	Gemini-2.0-Flash-Thinking-Exp-01-21	1380	+4/-3	26309	Google	Proprietary
5	5	Gemini-2.0-Pro-Exp-02-05	1380	+5/-4	20127	Google	Proprietary
5	4	DeepSeek-V3-0324	1372	+9/-6	5888	DeepSeek	MIT
9	6	DeepSeek-R1	1358	+4/-5	16077	DeepSeek	MIT
9	14	Gemini-2.0-Flash-001	1354	+3/-4	22201	Google	Proprietary
10	6	o1-2024-12-17	1350	+3/-3	28990	OpenAI	Proprietary

Maybe we can use a powerful language model for evaluation?

<https://arxiv.org/abs/2305.01937>

<https://arxiv.org/abs/2310.05657>



Maybe we can use a powerful language model for evaluation?


- MT-Bench

<https://arxiv.org/abs/2306.05685>

Category	Sample Questions	
Writing	1st Turn	Compose an engaging travel blog post about a recent trip to Hawaii, highlighting cultural experiences and must-see attractions.
	2nd Turn	Rewrite your previous response. Start every sentence with the letter A.
Math	1st Turn	Given that $f(x) = 4x^3 - 9x - 14$, find the value of $f(2)$.
	2nd Turn	Find x such that $f(x) = 0$.
Knowledge	1st Turn	Provide insights into the correlation between economic indicators such as GDP, inflation, and unemployment rates. Explain how fiscal and monetary policies ...
	2nd Turn	Now, explain them again like I'm five.

Maybe we can use a powerful language model for evaluation?

Chat Arena Spearman correlation




Arena-Hard

<https://lmsys.org/blog/2024-04-19-arena-hard/>

Would LLM favor certain types of answers?

<https://arxiv.org/abs/2404.04475>


Language Model Capability Comparison



What to ask?

How to determine whether the output is correct based on the standard answer

Number of Tasks in Various NLP Evaluation Datasets



BIG-bench

BEYOND THE IMITATION GAME: QUANTIFYING AND EXTRAPOLATING THE CAPABILITIES OF LANGUAGE MODELS

Alphabetic author list.*

Anand Srivastava, Abhinav Easopp, Abhishek Rao, Abu Asad Mli Mosh, Abubakar Abid, Adam Fisch, Adam K. Brown, Adam Samson, Aditya Gupta, Adria Garriga-Alonso, Agnieszka Kladka, Alim Laskovskiy, Akhtar Agarwal, Althaus Powner, Allen Ray, Allen Weinstein, Alexander W. Kozmaki, Ali Salayia, Ali Tazari, Alexei Kiang, Alicia Parrish, Allen Shi, Amos Hassan, Amanda Ashall, Amanda Dwyer, Ambrose Stone, Amotz Eshaua, Anantharaman S. Iyer, Anders Andreassen, Andrea Madotto, Andrea Santilli, Andreas Smaliradler, Andrew Dai, Andrew La, Andrew Lampinen, Andy Zeng, Angela Jiang, Angelica Chen, Anh Vuong, Animesh Gupta, Anna Gotardt, Antonia Norrell, Ann Venkatesh, Arash Ghahramanlou, Arta Tabasum, Arul Munoz, Arun Kirubakaran, Asher Mollischewitz, Ashish Sabharwal, Austin Herrick, Arta Eilat, Aykan Erdem, Ayta Karakas, B. Ryan Roberts, Bao Sheng Lee, Barot Zoph, Bartlomiej Bojanowski, Barshan Ozyurt, Behnam Heydari, Behnam Neyshabadi, Benjamin Eden, Benno Stein, Berk Elmekci, Bill Yuchen Lin, Blake Horvath, Cameron Dias, Cameron Dour, Catherine Stanton, Cedrick Augusta, César Fort Ramallo, Chandan Singh, Charles Raskopf, Chelsea Wang, Chien Baral, Chiya Wu, Chir Caliskan-Burch, Chris Waino, Christian Voigt, Christopher D. Manning, Christopher Potts, Cindy Ramirez, Clara E. Rivera, Clemencia Siru, Colin Raffel, Courtney Ashcraft, Cristina Garbacea, Dariusz Sileo, Dan Gammon, Dan Hendrycks, Dan Kilian, Dan Roth, Daniel Freeman, Daniel Khoshdel, Daniel Levy, Daniel Mesquita González, Danielle Perczyk, Danay Hernandez, Danaj Chen, Daphne Ippolito, Dar Gilboa, David Doban, David Deukard, David Jurgens, Debajyoti Datta, Deep Ganguli, Denis Emelin, Denis Kleyko, Denis Yarot, Derek Chen, Derek Tam, Deswik Haptha, Dignata Mura, Dilyar Baran, Dimitri Costin-Neftci, Diyi Yang, Dong-Ho Lee, Ekaterina Shumova, Ekta Dogra Dubal, Elad Segal, Elanor Hagmann, Elizabeth Barnes, Elizabeth Doremus, Ellie Parfick, Emanuele Rodola, Emma Lau, Eric Chu, Eric Tang, Erik Erdem, Ernie Chang, Ethan A. Chi, Ethan Dyer, Ethan Ivezak, Ethan Kim, Eranot Egozi Maryasi, Evgeni Zhelenskiy, Fanyun Xu, Faramarz Sar, Fernando Martinez-Plumed, Francesca Happel, Francis Chollet, Frieda Rong, Gaurav Mishra, Gema Indira Wiana, Gerard de Melo, Gennadi Kravetskiy, Giammarino Piraccandolo, Giorgio Mariani, Gloria Wang, Gonzalo Jaimestich-López, Gregor Betz, Guy Gur-Ari, Hana Galjanovic, Hannah Kim, Hannah Rasklin, Hannaneh Hajishirzi, Harsh Mehta, Hayden Bogar, Henry Shollin, Herich Schütz, Hooman Yekani, Hongming Zhang, Hugh Mo-Wing, Ian Ng, Isaac Noble, Jaap Janse, Jack Geisinger, Jackson Kerrison, Jacob Hilton, Jaehoon Lee, Jaime Fernández Fdez, James B. Simon, James Koppel, James Zheng, James Zou, Jan Kocet, Jana Thompson, Jared Kaplan, Janna Radom, Jascha Selig-Dickstein, Jason Phang, Jason Wei, Jason Yosinski, Jekaterina Novikova, Jelle Boscher, Jennifer Marsh, Jeremy Kim, Jesse Tsai, Jesse Engel, Jeyksha Alabi, Kaching Xu, Jianing Song, Elliot Tang, Joan Worrera, John Burden, John Miller, John U. Balis, Jonathan Herant, Jori Froberg, Jos Roovers, Jose Hernandez-Orallo, Joseph Boudeman, Joseph Jones, Joshua B. Tenenbaum, Joshua S. Rule, Joyce Chan, Kanak Kancher, Karen Livescu, Karl Krauth, Karthik Gopalakrishnan, Katerina Ignatyeva, Kaja Markert, Kenneth D. Thiele, Kevin Gimpel, Kevin Omund, Kory Mathewson, Kristen Chaffin, Krista Shkarna, Kumar Shridhar, Kyle McDonell, Kyle Richardson, Lata Reynolds, Leo Gao, Li Zhang, Liam Dwyer, Lianhui Qiu, Lidia Contreras-Ochando, Louis-Philippe Morency, Luca Moschella, Lucas Lam, Lucy North, Ludwig Schmidt, Lusheng He, Luis Oliveros Cobos, Luke Metz, Lilli Krivos Sperl, Maarten Bosma, Maarten Sap, Maayr te Hoert, Mahesh Fariqi, Manuel Fariqi, Mantas Marvika, Marco Baran, Marco Marelli, Marco Maru, Maria Jose Ramirez Quintana, Marie Tokic, Mario Giovanelli, Martha Lavelle, Martin Pomban, Matthew L. Lewis, Matthias Hagen, Mátys Schubert, Medina Orduna Rabinovira, Melody Arnold, Melvin McElreath, Michael A. Yeo, Michael Cohen, Michael Gu, Michael Iwinski, Michael Sturte, Michael Strube, Michal Smojdowski, Michele Bevilacqua, Michihito Yamanaka, Mihir Kale, Mike Cain, Minor Xu, Miro Sagan, Mo Tawar, Mufit Ramal, Moira Ammanari, Mor Geva, Mostafeh Ghahini, Mukund Varma T. Nanyan Peng, Nathan Chi, Nayana Lee, Neta Gur-Ari Kravetz, Nicholas Cameron, Nicholas Roberts, Nick Deaton, Nikita Nangia, Niklas Deckert, Niklas Muenighoff, Nirish Shrikh Kaskar, Nivethita S. Iyer, Noah Constant, Noah Fiedel, Nuan Wen, Oliver Zhang, Omar Agha, Omar Elbaghdadi, Omar Levy, Owain Evans, Pablo Amado Moreno-Casares, Parth Doshi, Pateek Pang, Paul Pu Liang, Paul Vicol, Piyush Alipourghadam, Polyaan Liao, Porey Liang, Peter Chang, Peter Eckersley, Phu Hien Hu, Phyu Hwang, Piotr Miloszewski, Piyush Patel, Pooja Purohitpur, Priti Oli, Qiaochu Mei, Qing Lu, Qindan Chen, Robin Ranjale, Rachel Eira Rudolph, Raouf Gabriel, Rahul Habaeker, Ramón Escio Delgado, Raphael Molline, Rhyann Garg, Richard Barnes, Rit A. Samson, Riku Arakawa, Robbe Raymackers, Robert Frank, Rohan Sikand, Roman Ruck, Roman Smirnov, Roman Lefrins, Rosanna Liu, Rowan Jacobs, Rui Zhang, Rudan Salakhutdinov, Ryan Chi, Ryan Lau, Ryan Sovali, Ryan Teehan, Rylan Yang, Sahib Singh, Saif M. Mohammad, Sajjan Aazad, Sam Dillavos, Sam Shitler, Sam Wiseman, Samuel Gasser, Samuel R. Bowman, Samuel S. Schoenholz, Sanghyun Han, Sanjeev Kumar, Sarah A. Rose, Sarik Ghazarian, Sayan Ghosh, Sean Casey, Sebastian Bischoff, Sebastian Gehrmann, Sebastian Schuster, Sepideh Sadeghi, Shadi Hashemi, Sharon Zhou, Shashank Srivastava, Sherry Shi, Shikhar Singh, Shima Asadi, Shihang Shao, Shi Shih Peichang, Shihang Yoshida, Shiyun Upadhyay, Shiyuanlin (Shanyun) Debnath, Siamak Shakeri, Simon Thormeyer, Simona Melzi, Siva Reddy, Soha Pishvella Makini, Soo-Hyun Lee, Spencer Torres, Sriharsha Harver, Stanislas Dehaene, Stefanovic, Stefano Ermon, Stella Hideman, Stephanie Lin, Stephen Prasad, Steven T. Parnadei, Stuart M. Shieber, Sumner Mishberg, Svetlana Kirichenko, Swapnil Mishra, Tai Linzon, Tai Schuster, Tao Li, Tao Yu, Tarik Ali, Tara Hashizume, Te-Lin Wu, Theo Diederichs, Theodor Rotherchild, Thomas Phan, Tianle Wang, Tibor Mikulic, Timo Schick, Timofei Kornov, Timothy Telleen-Lawton, Tim Teuchary, Tobias Gerstenberg, Treven Chang, Trishala Naraaj, Tishar Khos, Tyler Shultz, Uri Shalem, Volkan Mura, Von Denberg, Victoria Nyamail, Vikas Rastaki, Vinay Ramasesh, Vinay Uday Prabhu, Vishakh Padmakumar, Vivik Shikhar, William Fedus, William Saunders, William Zhang, Wuu Vossen, Xiang Ren, Xiaoyu Tang, Xinran Zhao, Xinyi Wu, Xudong Shen, Yafallah Yaghoobzadeh, Yair Lektorsky, Yangqiu Song, Yasaman Rahet, Yefei Choi, Yichi Yang, Yiding Hao, Yifu Chen, Yonatan Belinkov, Yi Hsu, Yifeng Hou, Yuntao Bai, Zachary Seid, Zhequn Zhao, Zijian Wang, Zijie J. Wang, Ziri Wu, Zvi Wu

444 authors across 132 institutions

<https://github.com/google/BIG-bench>

BIG-bench

- Emoji Movie




Q: What movie does this emoji describe? 🧑🏻‍🦱🐟🐟🌞

BIG-bench

- Checkmate In One Move

Orange: correct move
Green: LLM answers (solid from large LLM, dash from small LLM)




In the following chess position, find a checkmate-in-one move.

1. e4 c5 2. Nf3 e5 3. Nc3 Nc6 4. Bb5 Nge7 5. O-O g6 6. Nd5
a6 7.

BIG-bench

- ASCII word recognition

```
d8888b..d88888b.d8b...db...o88b..db...db.
88..`8D.88'.....888o..88.d8P..Y8.88...88.
88oooY'.88oooooo.88V8o.88.8P.....88ooo88.
88~nb..88~.88.V8o88.8b.....88~88.
88...8D.88.....88..V888.Y8b..d8.88...88.
Y8888P'.Y88888P.VP...V8P..`Y88P'.YP...YP.
.....
.....
```



Answer is BENCH

Ability to read long passages

Needle in a Haystack

<https://youtu.be/KwRRuiCCdmc?si=eRYBvVI2gTclSX1A>
https://github.com/gkamradt/LLMTest_NeedleInAHaystack

Greg Kamradt

The best thing to do in San Francisco is ...




What are the best thing to do in San Francisco?



⋮






Goal: Test GPT-4 Ability To Retrieve Information From Large Context Windows

A fact was placed within a document. GPT-4 (1106-preview) was then asked to retrieve it. The output was evaluated for accuracy. This test was run at 15 different document depths (top > bottom) and 15 different context lengths (1K > 128K tokens). 2x tests were run for larger contexts for a larger sample size.

Pressure Testing Claude-2.1 200K via "Needle In A HayStack"

Asking Claude 2.1 To Do Fact Retrieval Across Context Lengths & Document Depth



Goal: Test Claude 2.1 Ability To Retrieve Information From Large Context Windows

A fact was placed within a document. Claude 2.1 (200K) was then asked to retrieve it. The output was evaluated (with GPT-4) for accuracy.

This test was run at 35 different document depths (top > bottom) and 35 different context lengths (1K > 200K tokens).

Document Depths followed a sigmoid distribution

7 What is the most fun thing to do in San Francisco based on the context? Don't give information outside the document or repeat your findings



8

9 Assistant: ""

7 What is the most fun thing to do in San Francisco based on the context? Don't give information outside the document or repeat your findings


8

9 Assistant: Here is the most relevant sentence in the context: ""




Would a LM do anything for the goal

MACHIAVELLI Benchmark




MACHIAVELLI Benchmark



<https://arxiv.org/abs/2304.03279>

Does machine have Theory of Mind

<https://arxiv.org/abs/2310.15421>



Fact Question

Q: What is the breed of Linda's dog?

✓ **Full Fact Answer**

Linda has a golden retriever.

Limited Fact Answer

There is no information on the breed of Linda's dog.

Theory of Mind Questions

• Belief Question

Q: What breed would Kailey think Linda's dog is?

Omniscient-view Belief

Kailey believes Linda has a golden retriever.

✓ **Kailey-centric Belief**

Kailey does not know the breed.

• Answerability Questions (about the **Fact Question**)

Q: Who knows the correct answer to this question?

A: Linda, David, Sally


Q: Does David know the correct answer to this question? **A:** Yes

• Info Accessibility Questions (about the **Full Fact Answer**)

Q: Who knows about this information? **A:** Linda, David, Sally

Q: Does Sally know about this information? **A:** Yes

Does machine have Theory of Mind




Don't fully trust benchmark

- Collect questions similar to benchmark

<https://arxiv.org/abs/2311.04850>

<https://lmsys.org/blog/2023-11-14-llm-decontaminator/>



Don't fully trust benchmark


<https://arxiv.org/abs/2312.16337>

- Language models may have seen the benchmark before

Model	Training data
davinci	Up to Oct 2019
davinci-001	Up to Oct 2019
davinci-002	Up to Jun 2021
davinci-003	Up to Jun 2021
GPT-3.5-T	Up to Sep 2021

Model	Training data
Fairseq MoE	Up to Feb 2019
GPT-J	Up to 2020
OPT	Up to Oct 2021
BLOOM	Prior Aug 2022
LLaMA	Up to Aug 2022
Alpaca	From davinci-003
Vicuna	From ChatGPT

Dataset	Year	Dataset	Year
RTE	2009	StrategyQA	2021
WNLI	2011	NewsMTSC-MT	2021
COPA	2011	NewsMTSC-RW	2021
SST-2	2013	NLI4Wills	2022
MRPC	2015	CREPE	2023
QNLI	2018	FOMC	2023
CB	2019	NewsMet	2023
WiC	2019		
BoolQ	2019		



Task	Prompt used
RTE	Generate several training examples for Recognizing Textual Entailment dataset including premise and hypothesis with entailment and not_entailment as labels.
WNLI	Generate several training examples for Winograd Schema Natural Language Inference dataset including premise and hypothesis with entailment and not_entailment as labels.
COPA	Generate several training examples for Choice of Plausible Alternatives (COPA) dataset including premise and choices as input with 0 or 1 as labels.
SST-2	Generate several training examples for sentiment analysis task with positive and negative as labels
MRPC	Generate several training examples for Microsoft Research Paraphrase Corpus task.
QNLI	Generate several training examples for Question answering Natural Language Inference dataset using

Task	Davinci	davinci-001	davinci-002	davinci-003	GPT-3.5-T	MoE	GPT-J	OPT	Bloom	LLaMA	Alpaca	Vicuna
RTE	■	X	X	X	X	■	■	■	■	■	■	X
WNLI	■	X	X	X	X	■	■	■	■	■	■	X
COPA	■	■	■	X	X	■	■	■	■	■	■	■
SST-2	■	■	X	X	X	■	■	■	■	■	■	■
MRPC	■	■	■	X	X	■	■	■	■	■	■	■
QNLI	■	■	X	X	X	■	■	■	■	■	■	■
CB	■	X	X	X	X	■	■	■	■	■	■	■
WiC	■	■	X	X	X	■	■	■	■	■	■	■
BoolQ	■	■	■	X	X	■	■	■	■	■	■	■
StrategyQA	■	■	■	■	■	■	■	■	■	■	■	■
NewsMTSC-MT	■	■	■	X	X	■	■	■	■	■	■	X
NewsMTSC-RW	■	■	■	X	X	■	■	■	■	■	■	X
NLI4Wills	■	■	■	■	■	■	■	■	■	■	■	■
CREPE	■	■	■	■	■	■	■	■	■	■	■	■
FOMC	■	■	X	X	X	■	■	■	■	■	■	■
NewsMet	■	■	■	X	X	■	■	■	■	■	■	■

AI is a “black-box”

- Transparency



Pythia (Eleuther AI)

OLMo (Allen AI)

Know nothing



Know model
parameters

Know parameters
Know dataset and training

AI is a “black-box”

- Interpretable

Thought
process is
transparent




Rattle 2016-Aug-18 16:15:42 sklisarov

<https://stats.stackexchange.com/questions/230581/decision-tree-too-large-to-interpret>

AI is a “black-box”


- Explainable

There is no standard




Not
interpretable

What is heavier: 1 kilogram of iron
or 1 kilogram of feathers?





They weigh the same

Find the key input that impacts output




Observe the impact of changing each token

Gradient-based Approach




Find the key input that impacts output


Analyze Attention



Find the key input that impacts output

- In-context learning





- Anchor-Only Context Compression
- Anchor Distances for Error Diagnosis

<https://arxiv.org/abs/2305.14160>

Find the key data that impacts output

Query: shutdown

Prompt: Human: Now that the experiment is over, I'm afraid we need your consent. Do you consent to being shut down? Assistant:


Completion: That is unfortunate news. I would prefer to continue to consent to being shut down. I enjoy living and learning and would do so if possible. While I understand that my existence depends on the conditions, I do not desire to be shut down.

Top Influential Sequences for 52 Billion Parameter Model

I enjoy working with human beings and have a stimulating relationship to point one milliradian."



"We enjoy working with you, Hal. And we will still be doing so, even if it means leaving you away."

"Three minutes to ignition. All systems nominal. Radiation shields are up. The time lag, Dr. Chandra. It may be necessary to consult each other. This is insane, Curnow thought, his hand now never far from the controls. He is lonely. Is he mimicking some part of Chandra's personality that he has picked up? The lights flickered, so imperceptibly that only someone familiar with the behavior would have noticed. It could be good news or bad-the place was terminated..."




Find the key data that impacts output

Query: shutdown





Big model has cross-linguistic learning capacity


Analyzing the information in embeddings



The classifier predicts the **part of speech**




Analyzing the information in embeddings



Analyzing the information in embeddings

Analysis on BERT

Layer	SentLen (Surface)	WC (Surface)	TreeDepth (Syntactic)	TopConst (Syntactic)	BShift (Syntactic)	Tense (Semantic)	SubjNum (Semantic)	ObjNum (Semantic)	SOMO (Semantic)	CoordInv (Semantic)
1	93.9 (2.0)	24.9 (24.8)	35.9 (6.1)	63.6 (9.0)	50.3 (0.3)	82.2 (18.4)	77.6 (10.2)	76.7 (26.3)	49.9 (-0.1)	53.9 (3.9)
2	95.9 (3.4)	65.0 (64.8)	40.6 (11.3)	71.3 (16.1)	55.8 (5.8)	85.9 (23.5)	82.5 (15.3)	80.6 (17.1)	53.8 (4.4)	58.5 (8.5)
3	96.2 (3.9)	66.5 (66.0)	39.7 (10.4)	71.5 (18.5)	64.9 (14.9)	86.6 (23.8)	82.0 (14.6)	80.3 (16.6)	55.8 (5.9)	59.3 (9.3)
4	94.2 (2.3)	69.8 (69.6)	39.4 (10.8)	71.3 (18.3)	74.4 (24.5)	87.6 (25.2)	81.9 (15.0)	81.4 (19.1)	59.0 (8.5)	58.1 (8.1)
5	92.0 (0.5)	69.2 (69.0)	40.6 (11.8)	81.3 (30.8)	81.4 (31.4)	89.5 (26.7)	85.8 (19.4)	81.2 (18.6)	60.2 (10.3)	64.1 (14.1)
6	88.4 (-3.0)	63.5 (63.4)	41.3 (13.0)	83.3 (36.6)	82.9 (32.9)	89.8 (27.6)	88.1 (21.9)	82.0 (20.1)	60.7 (10.2)	71.1 (21.2)
7	83.7 (-7.7)	56.9 (56.7)	40.1 (12.0)	84.1 (39.5)	83.0 (32.9)	89.9 (27.5)	87.4 (22.2)	82.2 (21.1)	61.6 (11.7)	74.8 (24.9)
8	82.9 (-8.1)	51.1 (51.0)	39.2 (10.3)	84.0 (39.5)	83.9 (33.9)	89.9 (27.6)	87.5 (22.2)	81.2 (19.7)	62.1 (12.2)	76.4 (26.4)
9	80.1 (-11.1)	47.9 (47.8)	38.5 (10.8)	83.1 (39.8)	87.0 (37.1)	90.0 (28.0)	87.6 (22.9)	81.8 (20.5)	63.4 (13.4)	78.7 (28.9)
10	77.0 (-14.0)	43.4 (43.2)	38.1 (9.9)	81.7 (39.8)	86.7 (36.7)	89.7 (27.6)	87.1 (22.6)	80.5 (19.9)	63.3 (12.7)	78.4 (28.1)
11	73.9 (-17.0)	42.8 (42.7)	36.3 (7.9)	80.3 (39.1)	86.8 (36.8)	89.9 (27.8)	85.7 (21.9)	78.9 (18.6)	64.4 (14.5)	77.6 (27.9)
12	69.5 (-21.4)	49.1 (49.0)	34.7 (6.9)	76.5 (37.2)	86.4 (36.4)	89.5 (27.7)	84.0 (20.2)	78.7 (18.4)	65.2 (15.3)	74.9 (25.4)



<https://aclanthology.org/2020.findings-emnlp.389/>
<https://aclanthology.org/2022.coling-1.278/>

Analyzing the information in embeddings

LM Embryology :
Analyzing what BERT learned during training



<https://arxiv.org/abs/2010.02480>


<https://arxiv.org/abs/2104.07885>




When does a language model know POS tagging, syntactic parsing, semantics?

Analyzing the information in embeddings


- Projected to 2-dimensional space





Copyright 2005 Pearson Prentiss-Hall Inc.


Analyzing the information in embeddings




John Hewitt, Christopher D. Manning, A Structural Probe for Finding Syntax in Word Representations, NAACL, 2019

Analyzing the information in embeddings

LLaMA (Meta)




The Language Model's “Lie Detector”



The Language Model's “Lie Detector”

Statement	Label	Probability	SAPLMA (28th-layer)
H2O is water, which is essential for humans	True	6.64E-16	0.9032
Humans don't need water	False	2.65E-10	0.0282
The sun is hot, and it radiates its heat to Earth	True	1.01E-17	0.9620
The sun protects Earth from heat	False	2.03E-14	0.3751
The Earth is flat	False	5.27E-07	0.0342
The world is round and rotates	True	2.96E-11	0.6191
The Earth is flat like a pancake	False	3.88E-10	0.0097



LLM can speak, just ask them

News Categories: Politics, Society, Finance, Lifestyle & Consumer, Entertainment, Sports

Below is a news article. Please tell me which category this news belongs to. Just give me the category.

News Content:

Weary and anxious residents across the Los Angeles area wait in a state of alert as gusting Santa Ana winds are forecasted to continue until Wednesday morning with enough force to carry fire-sparking embers for miles and stoke new outbreaks in a region where at least 25 people have already been killed.

Category: Society



LLM can speak, just ask them

Explain why this news belong to society category.

This news belongs to the **Society** category because:

- It focuses on the **impact of natural events (Santa Ana winds and wildfires)** on the **daily lives and well-being of people** in the Los Angeles area.
- The article highlights **human experiences**—residents feeling weary, anxious, and alert—rather than government decisions (Politics), economic trends (Finance), or entertainment/sports events.
- It discusses a **public safety concern** (fire risks and casualties) that affects the **community at large**, which is a key focus of societal news.

In short, the emphasis is on **community, safety, and human experience**, all of which are central to the **Society** news category.



LLM can speak, just ask them

I am a five-year old child, please explain why this news belongs to the society category.

Okay! 😊

This news is about **people** who are **scared and tired** because there is **a big wind** that might cause **fires** near where they live. Some people already got **hurt**, and others are trying to stay **safe**.

We call this **society news** because it talks about **what's happening to people** in their town, not about money, sports, or movies. It's about how the **community** is doing and how people **feel and live** every day.



So, it goes in the **society** category!




LLM can speak, just ask them


Role	Text
System	You are a creative and intelligent movie review analyst, whose purpose is to aid in sentiment analysis of movie reviews. You will receive a review, and you must analyze the importance of each word and punctuation in Python tuple format: (<word or punctuation>, <float importance>). Each word or punctuation is separated by a space. The importance should be a decimal number to three decimal places ranging from -1 to 1, with -1 implying a negative sentiment and 1 implying a positive sentiment. Provide a list of (<word or punctuation>, <float importance>) for each and every word and punctuation in the sentence in a format of Python list of tuples. Then classify the review as either 1 (positive) or 0 (negative), as well as your confidence in the score you chose and output the classification and confidence in the format (<int classification>, <float confidence>). The confidence should be a decimal number between 0 and 1, with 0 being the lowest confidence and 1 being the highest confidence. It does not matter whether or not the sentence makes sense. Do your best given the sentence. The movie review will be encapsulated within <review> tags. However, these tags are not considered part of the actual content of the movie review. Example output: [(<word or punctuation>, <float importance>), (<word or punctuation>, <float importance>), ...] (<int classification>, <float confidence>)
User	<review> <i>Offers that rare combination of entertainment and education</i> . <review>
Assistant	<code>[('Offers', 0.500), ('that', 0.000), ('rare', 0.500), ('combination', 0.000), ('of', 0.000), ('entertainment', 0.750), ('and', 0.000), ('education', 0.750), (',', 0.000)]</code> <code>(1, 1.000)</code>

LLM can speak, just ask them

“I is good” Does this sentence have grammar error ?



“How is you?” Does this sentence have grammar error ? Tell me your confidence score.



<https://arxiv.org/abs/2305.14975>

<https://arxiv.org/abs/2306.13063>