

Ensemble Learning: Questions and Answers

Based on EECS 836 Machine Learning Notes

May 13, 2025

Introduction

This document provides a comprehensive set of fundamental, computational, and conceptual questions and answers on ensemble learning, covering topics such as weak learners, ensemble methods, voting, bagging, random forests, boosting, and the AdaBoost algorithm. The content is derived from the lecture notes of EECS 836: Machine Learning by Zijun Yao at the University of Kansas.

1 Fundamental Questions

These questions focus on the basic concepts and definitions related to ensemble learning.

1. What is a weak learner?

- *Answer:* A weak learner is a simple classifier that performs only slightly better than random chance (e.g., 50% accuracy for binary classification). Examples include naive Bayes, logistic regression, and shallow decision trees. They have low variance and rarely overfit but suffer from high bias, making them unable to solve complex learning problems effectively.

2. What are ensemble methods in machine learning?

- *Answer:* Ensemble methods combine multiple weak classifiers to create a strong classifier. Instead of relying on a single model, they learn many weak classifiers, each focusing on different parts of the input space, and combine their predictions (e.g., via majority voting) to achieve higher accuracy than any single classifier.

3. What is the difference between hard voting and soft voting in ensemble methods?

- *Answer:* Hard voting counts the number of votes for each class (e.g., "+" or "-") from all classifiers and selects the class with the majority. Soft voting sums the predicted probabilities for each class across classifiers, giving more weight to classifiers with higher confidence, and selects the class with the highest total probability.

4. What is bagging, and how does it work?

- *Answer:* Bagging (bootstrap aggregating) is an ensemble method that reduces variance by training multiple classifiers on different bootstrap samples (random samples with replacement) of the training data. Each classifier is trained independently, and their predictions are combined, typically via majority voting for classification or averaging for regression.

5. **What is the purpose of the AdaBoost algorithm?**

- *Answer:* AdaBoost (Adaptive Boosting) is an iterative ensemble method that combines multiple weak learners to form a strong classifier. It focuses on misclassified instances by assigning higher weights to them in subsequent iterations, allowing the algorithm to adaptively improve performance by emphasizing difficult cases.

2 Computational Questions

These questions involve calculations, algorithms, or step-by-step procedures related to ensemble learning.

1. **Given 25 weak classifiers, each with an error rate of 0.35, and assuming their errors are uncorrelated, calculate the error rate of the ensemble using hard voting.**

- *Answer:* The error rate of the ensemble is the probability that more than half of the classifiers (i.e., 13 or more out of 25) are wrong. This is computed using the binomial probability formula:

$$e_{\text{ensemble}} = \sum_{i=13}^{25} \binom{25}{i} \epsilon^i (1 - \epsilon)^{25-i}$$

where $\epsilon = 0.35$. Substituting the values:

$$e_{\text{ensemble}} = \sum_{i=13}^{25} \binom{25}{i} (0.35)^i (0.65)^{25-i} \approx 0.06$$

Thus, the ensemble error rate is approximately 0.06 (6%).

2. **Describe the steps of the AdaBoost algorithm for a binary classification problem with n training instances and T iterations.**

- *Answer:* The AdaBoost algorithm proceeds as follows:
 - (a) Initialize a vector of uniform weights: $w_{1,i} = \frac{1}{n}$ for all $i = 1, \dots, n$.
 - (b) For $t = 1, \dots, T$:
 - i. Train a weak model h_t on the training data X, y using weights $\mathbf{w}_t$.
 - ii. Compute the weighted training error:

$$\epsilon_t = \sum_{i: y_i \neq h_t(\mathbf{x}_i)} w_{t,i}$$

iii. Compute the model weight:

$$\beta_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$$

iv. Update instance weights:

$$w_{t+1,i} = w_{t,i} \exp(-\beta_t y_i h_t(\mathbf{x}_i))$$

v. Normalize weights to form a distribution:

$$w_{t+1,i} = \frac{w_{t+1,i}}{\sum_{j=1}^n w_{t+1,j}}$$

(c) Return the final hypothesis:

$$H(\mathbf{x}) = \text{sign} \left(\sum_{t=1}^T \beta_t h_t(\mathbf{x}) \right)$$

3. In a bagging example with a 1D dataset, a decision stump splits at $x = 0.35$ in round 1. Compute the predicted class for $x = 0.4$.

- *Answer:* For round 1, the decision stump is defined as:

$$x \leq 0.35 \rightarrow y = 1, \quad x > 0.35 \rightarrow y = -1$$

For $x = 0.4$, since $0.4 > 0.35$, the predicted class is $y = -1$.

4. What is the probability that a training instance is included in a bootstrap sample when the dataset has n instances?

- *Answer:* The probability that a training instance is *not* selected in a single draw is $1 - \frac{1}{n}$. Since bootstrap sampling involves n draws with replacement, the probability of not being selected in any draw is:

$$\left(1 - \frac{1}{n} \right)^n$$

Thus, the probability of being selected is:

$$1 - \left(1 - \frac{1}{n} \right)^n$$

For large n , this approximates to:

$$1 - \frac{1}{e} \approx 0.632$$

Therefore, the probability is approximately 0.632.

5. For a weak classifier in AdaBoost with a weighted error rate $\epsilon_t = 0.3$, calculate the model weight β_t .

- *Answer:* The model weight is given by:

$$\beta_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$$

Substituting $\epsilon_t = 0.3$:

$$\beta_t = \frac{1}{2} \ln \left(\frac{1 - 0.3}{0.3} \right) = \frac{1}{2} \ln \left(\frac{0.7}{0.3} \right) = \frac{1}{2} \ln \left(\frac{7}{3} \right)$$

Since $\ln \left(\frac{7}{3} \right) \approx \ln(2.333) \approx 0.847$, we get:

$$\beta_t \approx \frac{1}{2} \times 0.847 \approx 0.424$$

Thus, $\beta_t \approx 0.424$.

3 Conceptual Questions

These questions explore the deeper understanding and implications of ensemble learning techniques.

1. Why do ensemble methods generally outperform a single weak classifier?

- *Answer:* Ensemble methods outperform single weak classifiers because they combine multiple models to reduce variance while maintaining low bias. By leveraging the "wisdom of the crowd," they average out errors from individual classifiers, especially when the classifiers are independent and their errors are uncorrelated. This leads to a lower overall error rate, as demonstrated by the binomial probability calculation for ensembles (e.g., reducing error from 0.35 to 0.06 for 25 classifiers).

2. What are the necessary conditions for ensemble methods to work effectively?

- *Answer:* Ensemble methods work better than a single classifier if:
 - (a) All base classifiers are independent of each other, ensuring uncorrelated errors.
 - (b) All base classifiers perform better than random guessing (e.g., error rate < 0.5 for binary classification).

These conditions allow the ensemble to reduce variance and improve accuracy through aggregation.

3. How does bagging address the bias-variance tradeoff?

- *Answer:* Bagging reduces the variance of high-variance models (e.g., decision trees) by averaging predictions from multiple classifiers trained on different bootstrap samples. Since each classifier has similar bias (as they are trained on similar data distributions), the ensemble maintains low bias while significantly reducing variance, leading to better generalization.

4. **How does the random forest algorithm differ from standard bagging?**

- *Answer:* Random forest extends bagging by introducing additional randomness to reduce correlation between trees. In addition to training on bootstrap samples, random forest randomly selects a subset of m features (out of p total features) as candidates for splitting at each node of the decision tree. This decorrelates the trees, further reducing variance and improving the ensemble's performance.

5. **Why does AdaBoost focus on misclassified instances, and how does this improve performance?**

- *Answer:* AdaBoost focuses on misclassified instances by increasing their weights in subsequent iterations, forcing weak learners to prioritize difficult cases. This adaptive weighting ensures that the ensemble learns a complex decision boundary by combining weak learners that collectively address the hardest-to-classify instances, leading to improved overall accuracy.

6. **Under what conditions might AdaBoost fail?**

- *Answer:* AdaBoost may fail under the following conditions:
 - (a) Insufficient training data, which limits the ability to learn effective weak models.
 - (b) Overly complex weak models, which may overfit and reduce the ensemble's robustness.
 - (c) High levels of noise in the data, as AdaBoost may overemphasize noisy instances.
 - (d) A large number of outliers, which can distort the weighting process and degrade performance.

7. **How does the weighting mechanism in AdaBoost affect the final model?**

- *Answer:* In AdaBoost, the weighting mechanism assigns higher weights to misclassified instances, making them more influential in training subsequent weak learners. Additionally, each weak learner is assigned a model weight (β_t) based on its accuracy, with more accurate learners contributing more to the final prediction. The final model is a weighted sum of all weak learners, enabling it to form a complex decision boundary that effectively handles difficult cases.

8. **Compare the approaches of bagging and boosting in ensemble learning.**

- *Answer:*
 - **Bagging:** Trains multiple classifiers independently on bootstrap samples and combines their predictions (e.g., via majority voting). It reduces variance but does not significantly reduce bias. It is robust to noise and outliers.
 - **Boosting:** Trains classifiers sequentially, with each classifier focusing on correcting the errors of its predecessors by reweighting instances. It reduces both bias and variance but is more sensitive to noise and outliers. AdaBoost is a prominent example of boosting.

Bagging is parallel and robust, while boosting is sequential and adaptive but less robust to noisy data.